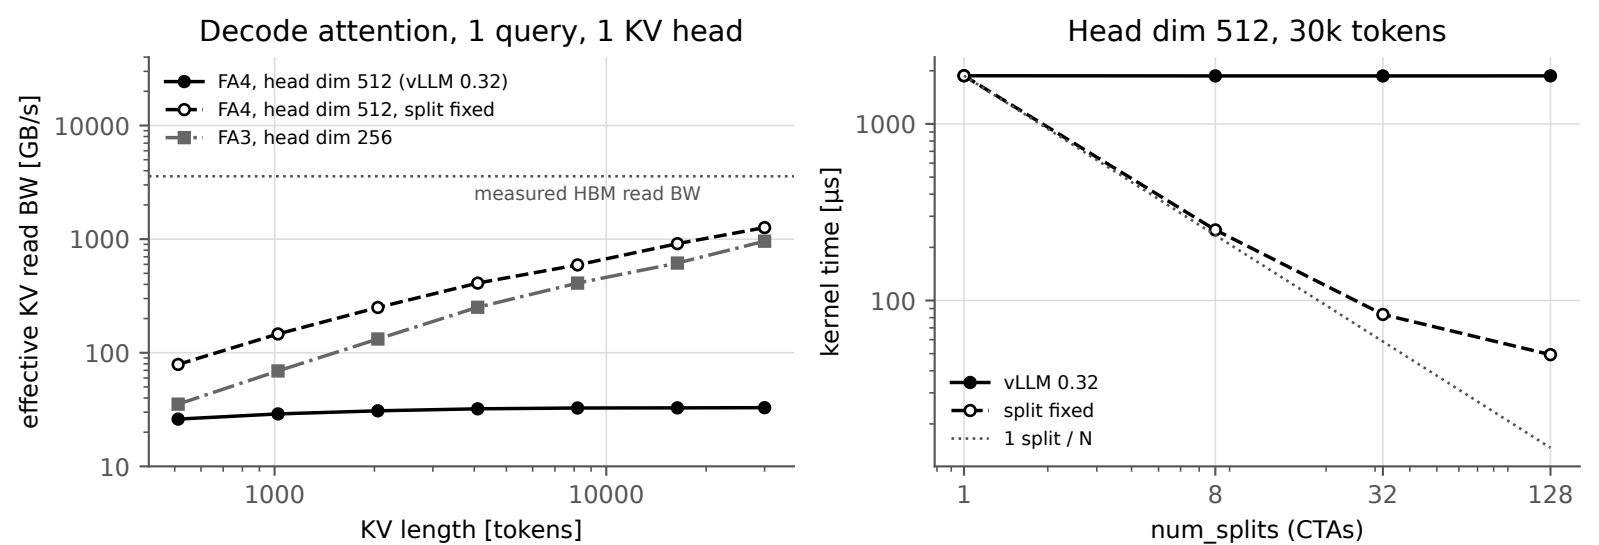
Decode attention, 1 query, 1 KV head
FA4, head dim 512 (vLLM 0.32)
FA4, head dim 512, split fixed
FA3, head dim 256
measured HBM read BW
effective KV read BW [GB/s]
10000
1000
100
10
1000
10000
KV length [tokens]
Head dim 512, 30k tokens
kernel time [µs]
1000
100
vLLM 0.32
split fixed
1 split / N
1
8
32
128
num_splits (CTAs)