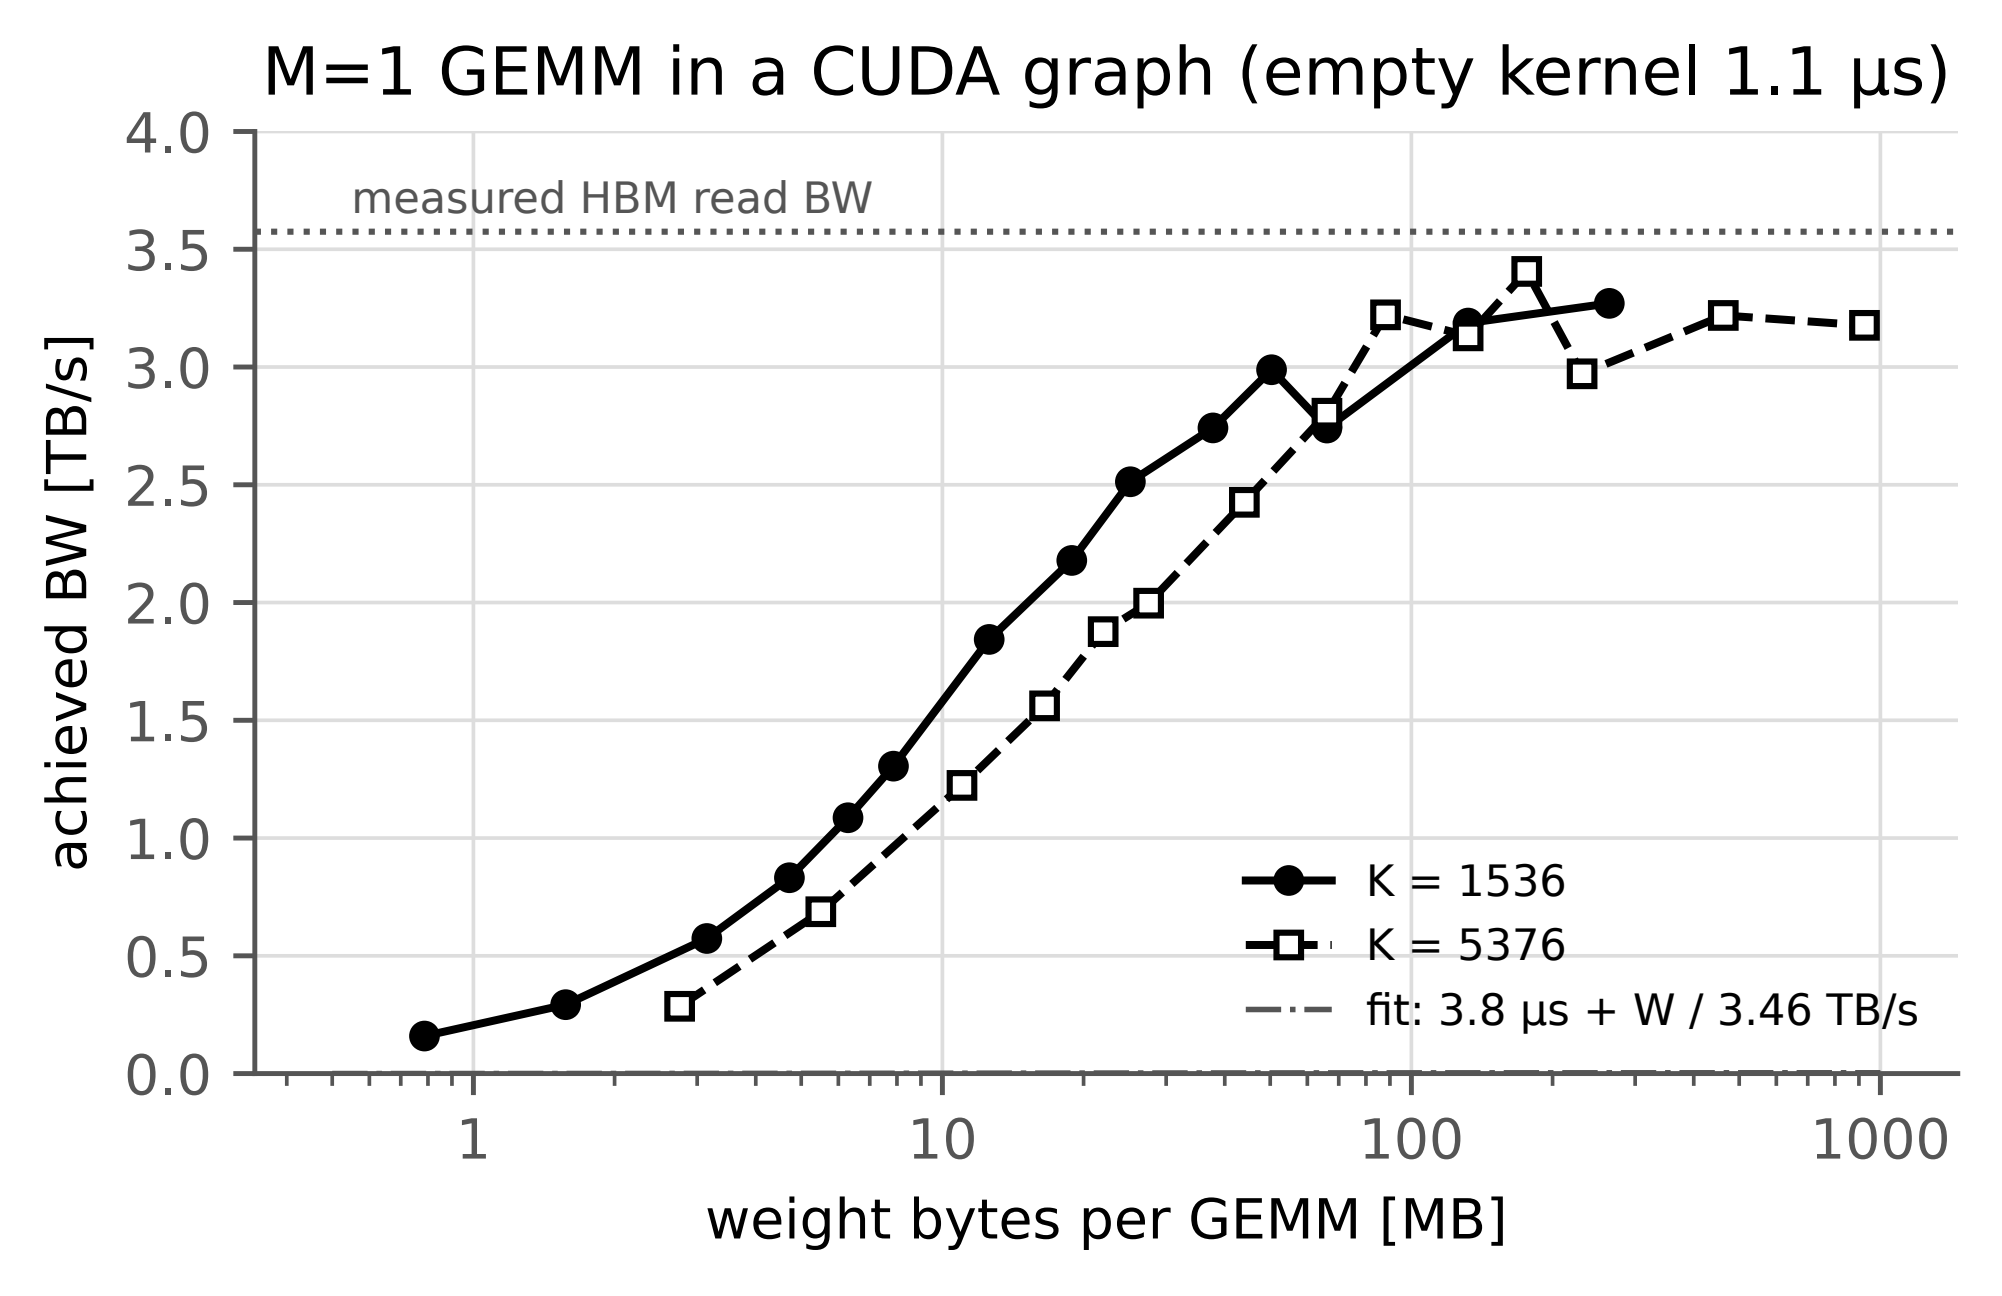
M=1 GEMM in a CUDA graph (empty kernel 1.1 µs)
measured HBM read BW
achieved BW [TB/s]
weight bytes per GEMM [MB]
K = 1536
K = 5376
fit: 3.8 µs + W / 3.46 TB/s
0.0
0.5
1.0
1.5
2.0
2.5
3.0
3.5
4.0
1
10
100
1000